SDLE

Local-First Collaborative Shopping List

T07 - G13

Henrique Silva

João Fernandes

Tomás Palma

Architecture Overview

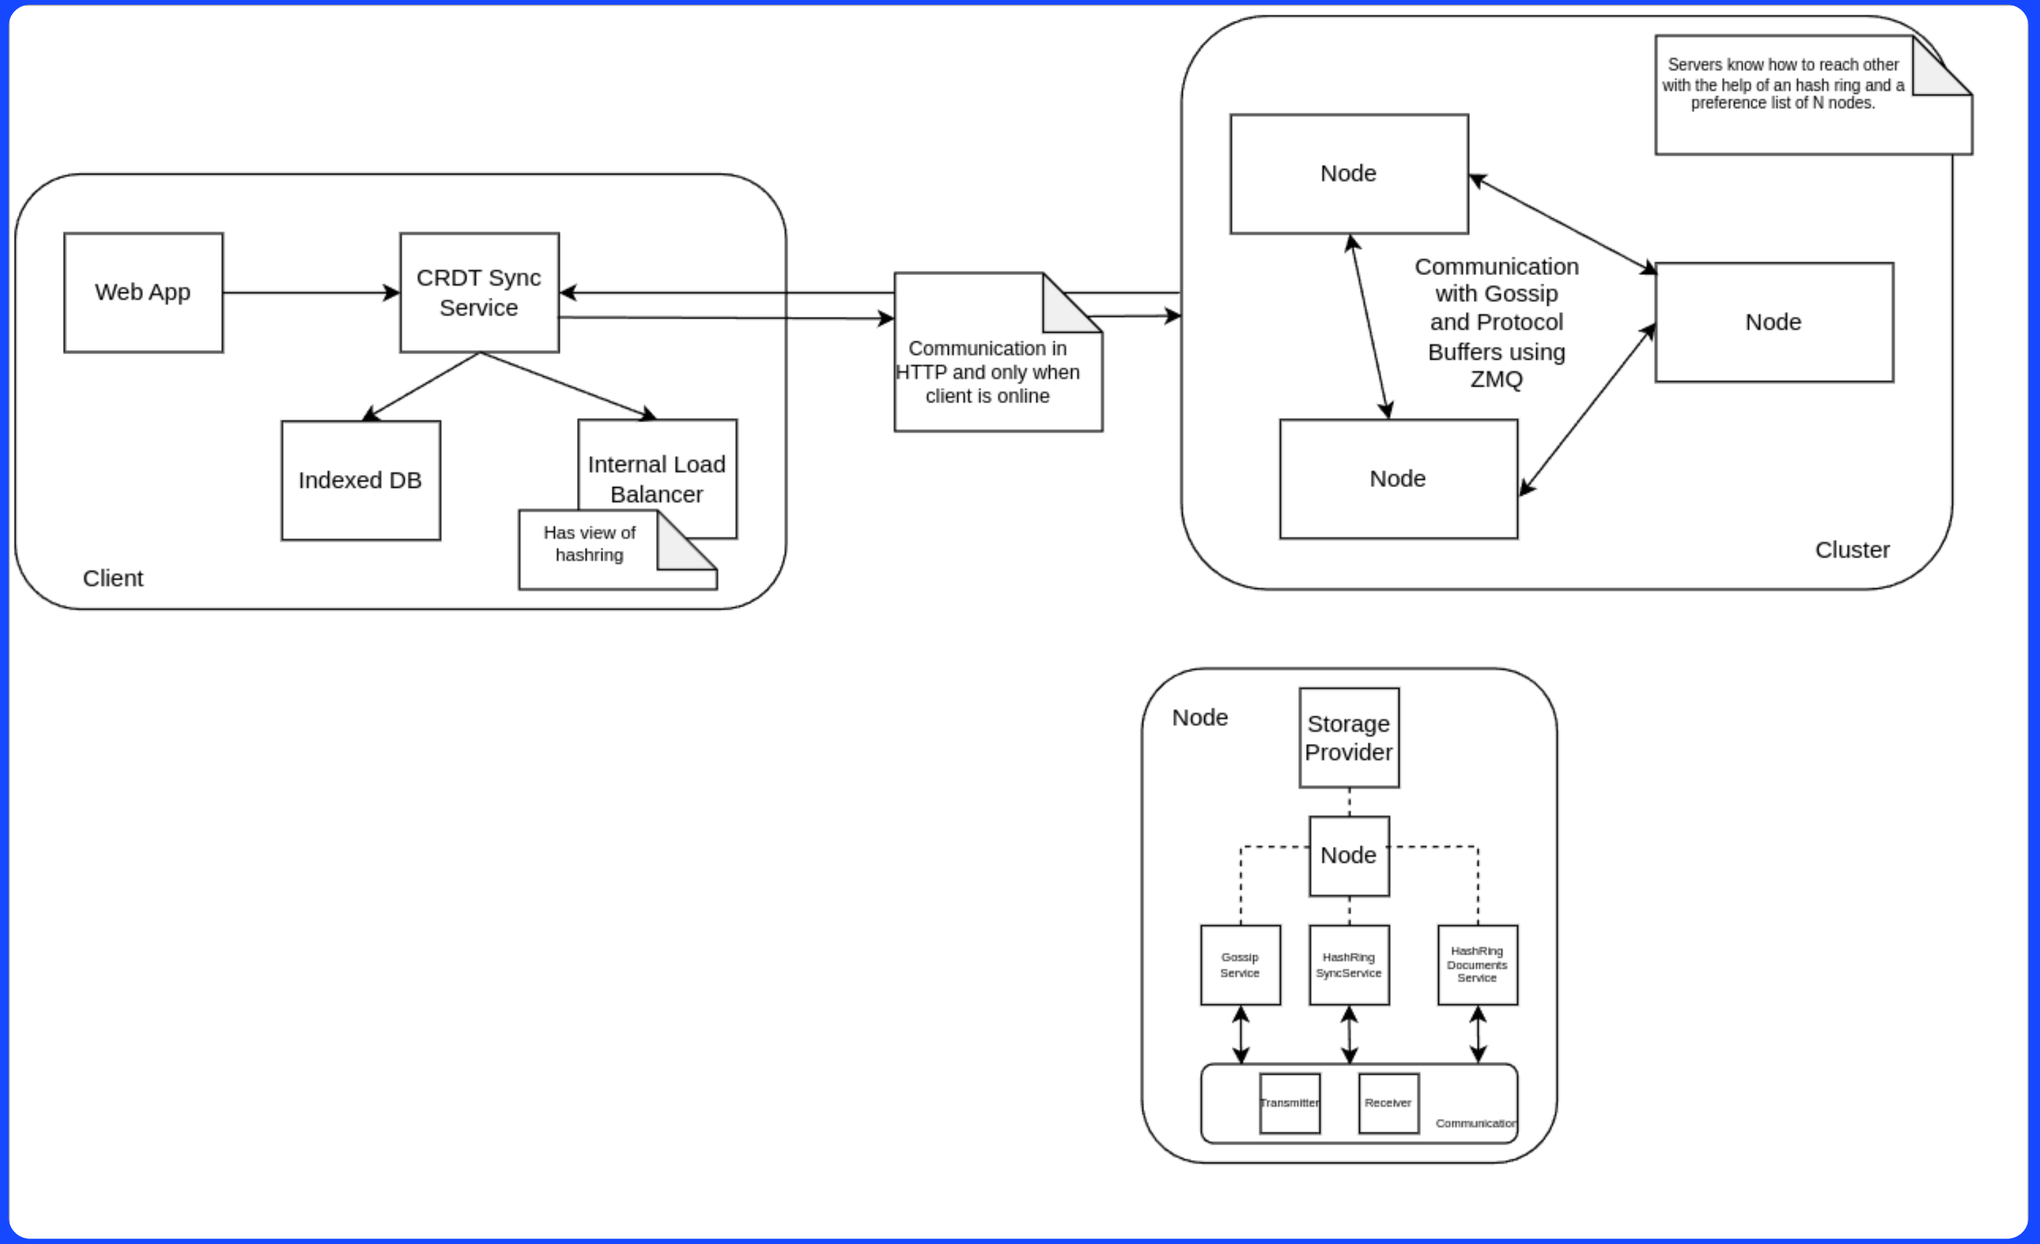

Serialization and Deserialization

Frontend to Backend

{JSON}

- Better support with javascript
- Less planned work to implement

Backend (Between server nodes)

protobuf
Protocol Buffers

- Binary format with better performance

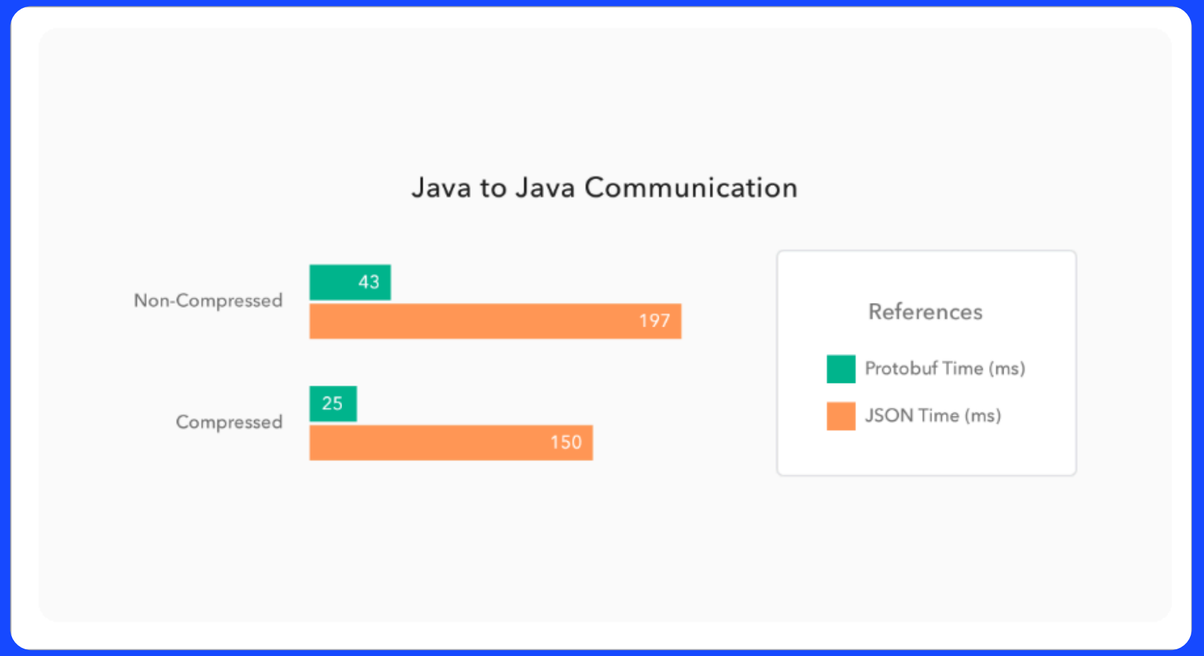

Chart generated by Okta with the performance of
500 GET requests between spring boot applications

Client-Side Load Balancing (1 /2)

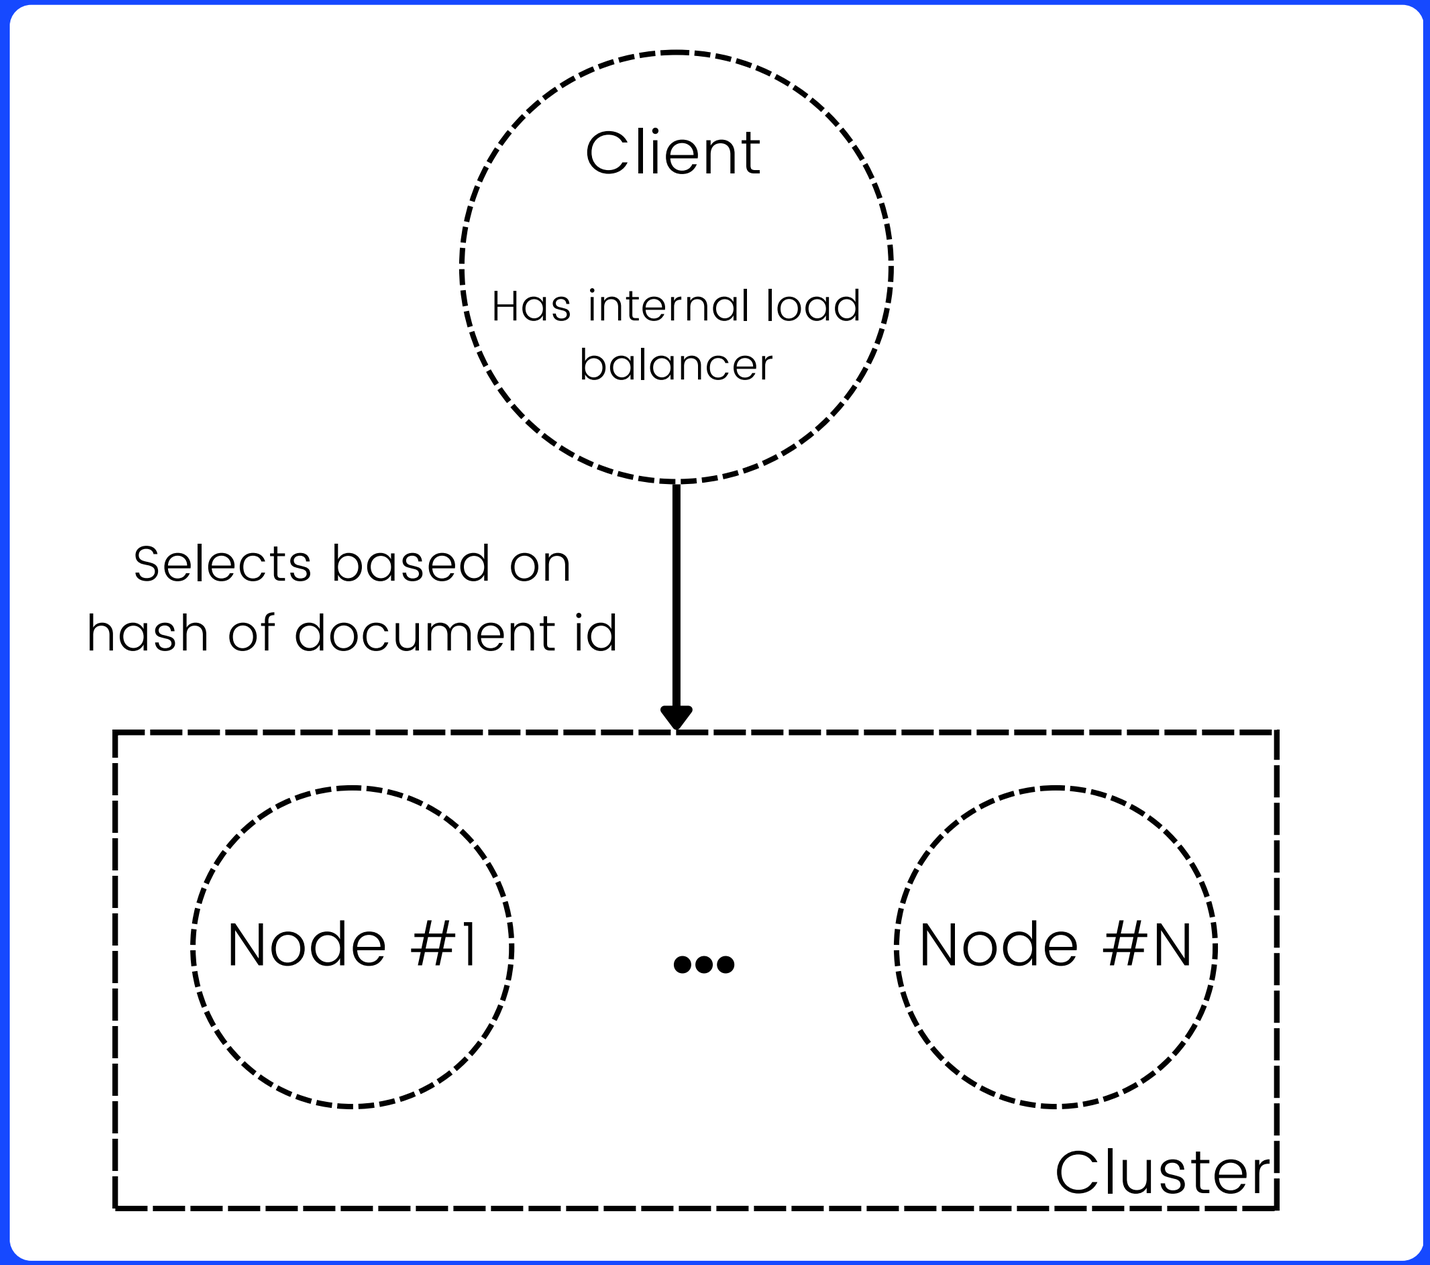

Client-Side Load Balancing (2/2)

- Reduces latency as it reduces round hops
- Not necessarily insecure since the client is controlled by the web server serving the web app

Table 2: Performance of client-driven and server-driven coordination approaches.

	99.9th percentile read latency (ms)	99.9th percentile write latency (ms)	Average read latency (ms)	Average write latency (ms)
Server-driven	68.9	68.5	3.9	4.02
Client-driven	30.4	30.4	1.55	1.9

Dynamo Paper

Ring partioning scheme (1 /2) - Strategies Explored

STRATEGY NAME	FUNCTIONING
STRATEGY 1	Assign randomly the hash of a node to the ring
STRATEGY 2	Assign random hashes to nodes and decide which node to take which key based on well defined partitions
STRAGEGY 3	Divide the hash ring into well defined partitions and then assign partitions to nodes

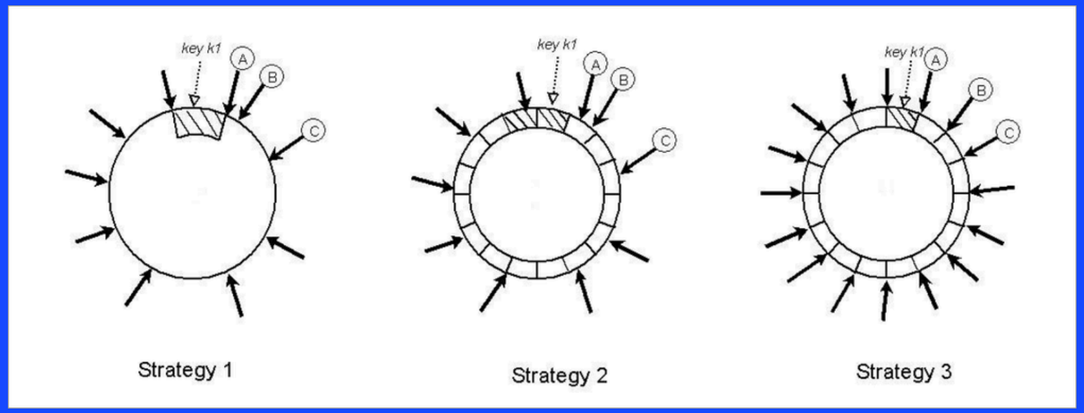

Dynamo Paper

Ring partioning scheme (2 /2) - Decision

- Strategy 2 was discarded
- Strategy 1 and Strategy 3 do not have significant efficiency differences
- Since strategy 1 was simpler to develop, that was the chosen one
- The hashring is modeled as a TreeMap<BigInteger, NodeIdentifier> where each node has more than one entry

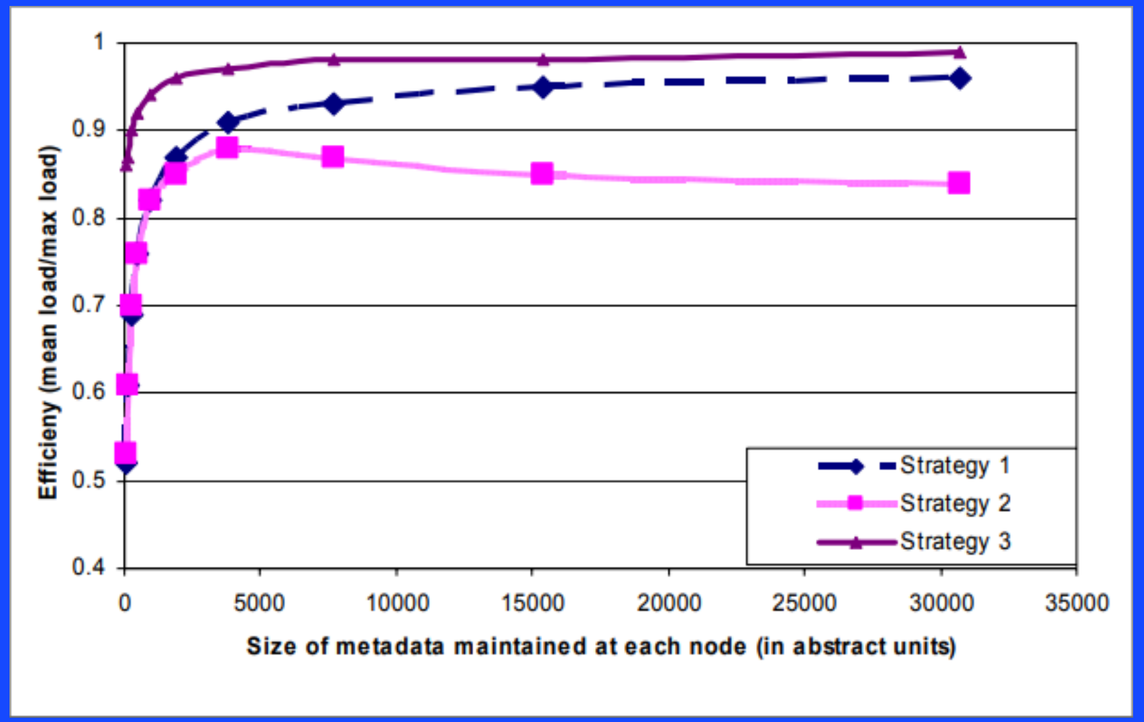

Sloppy quorums & Hinted Handoffs

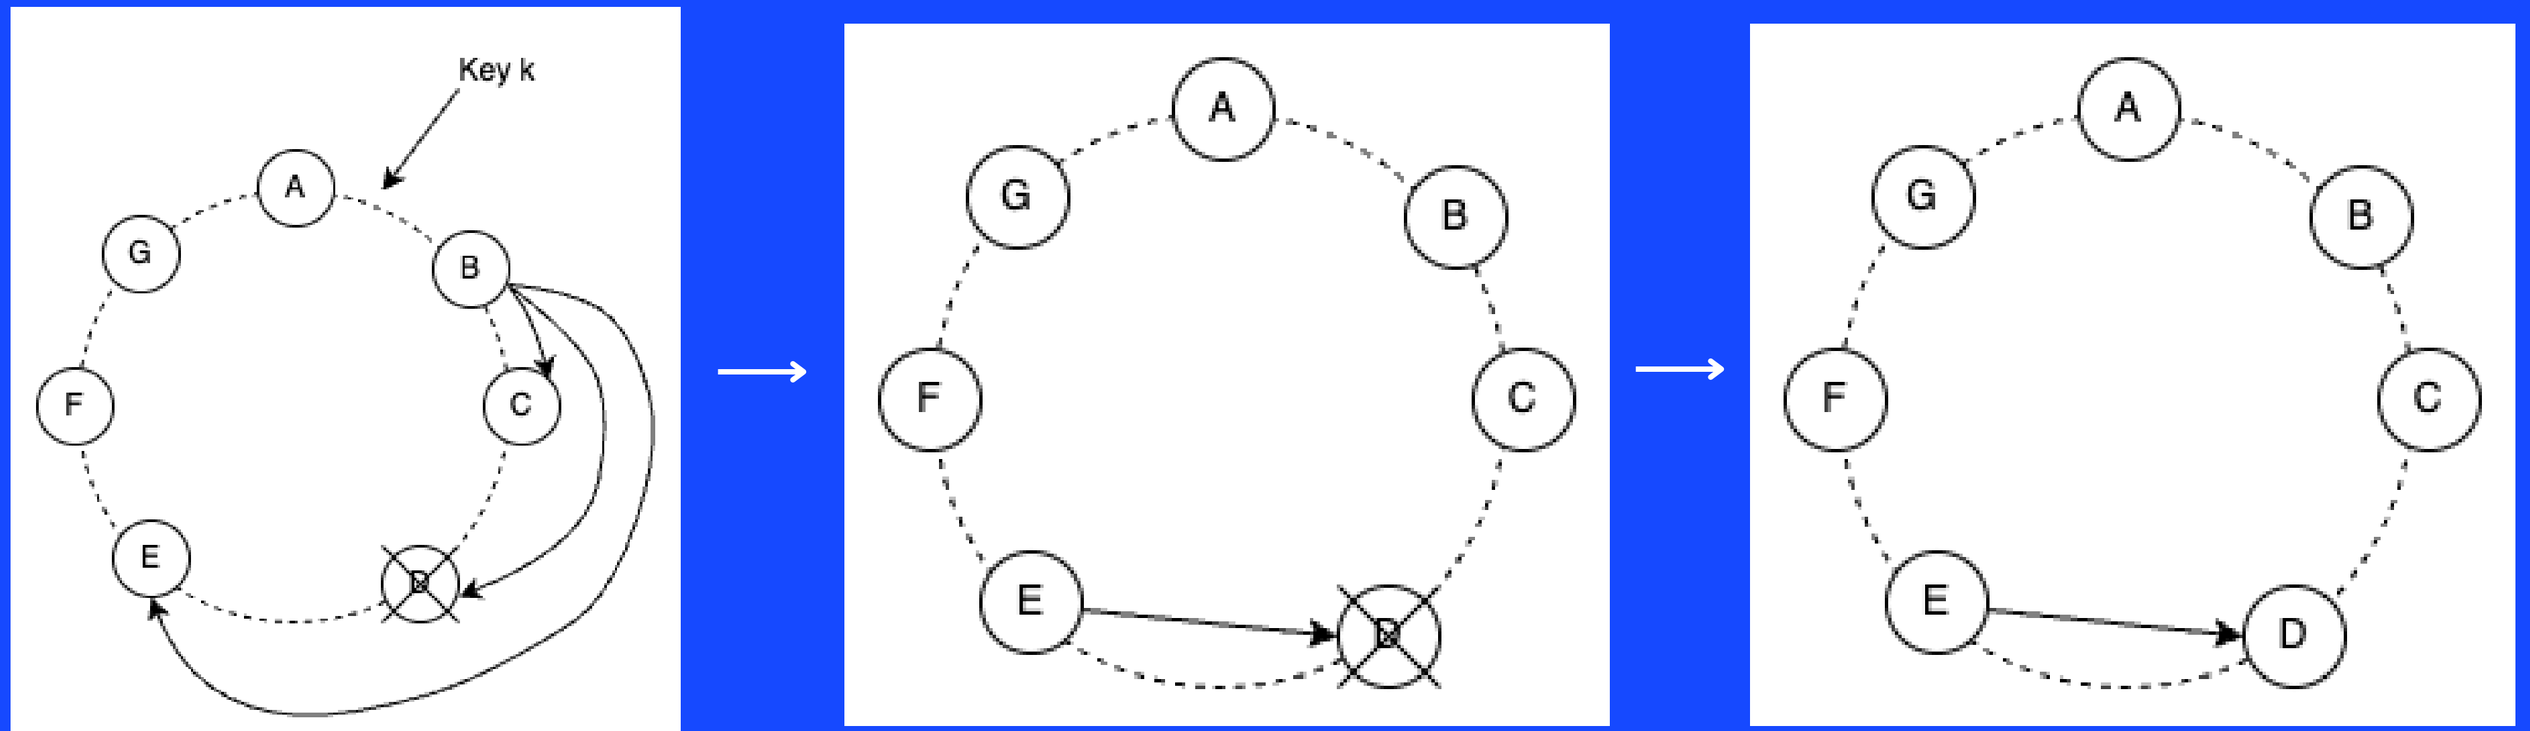

Reason neither pub-sub nor proxy patterns were used (1/2)

1. Single point of failure

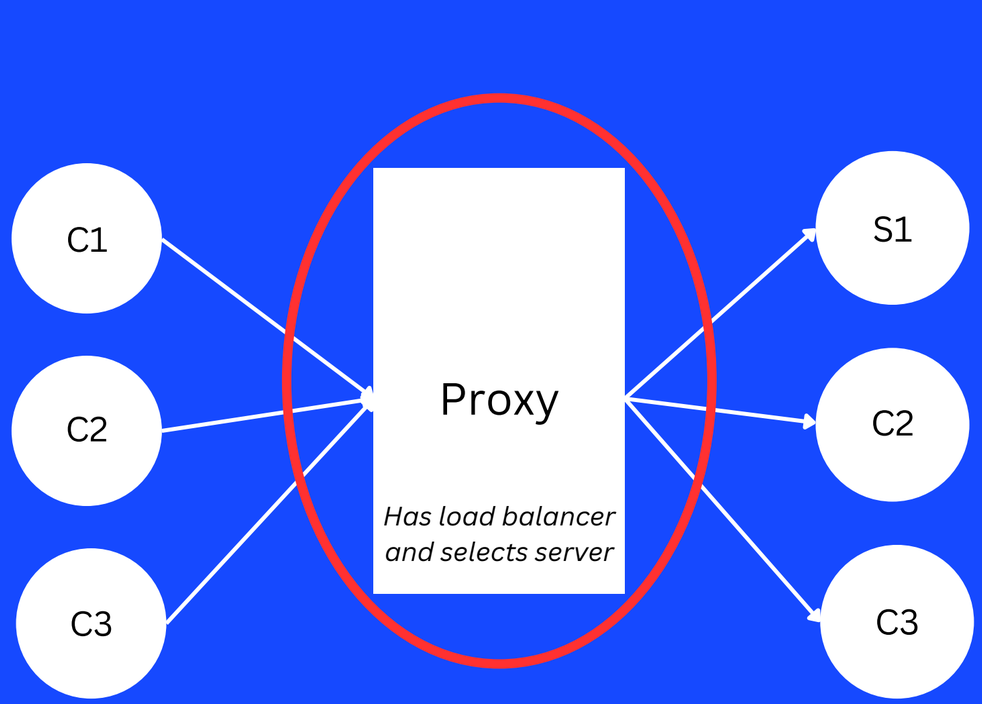

Proxy pattern

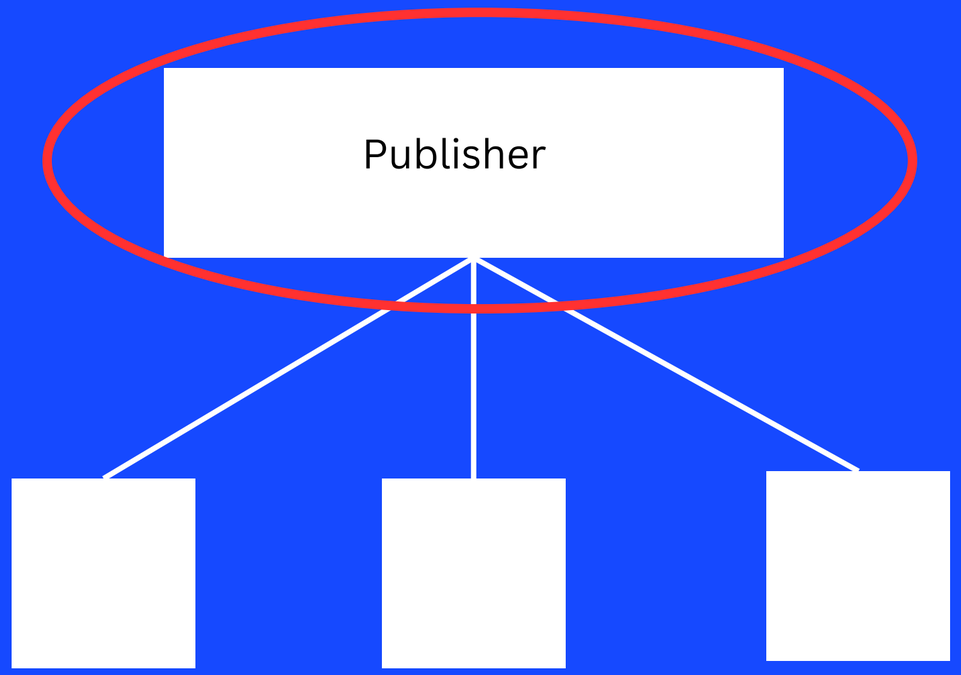

Pub-Sub pattern

Although fixed by DNS, it remains a centralized approach

Reason neither pub-sub nor proxy patterns were used (2/2)

2. Other notes

- Even if we had many servers under DNS we would still have to have a conflict resolution mechanism, so we chose to try to implement one.

- The assignment of IP Addresses of a certain name is not decentralized

Log-CRDT

- Used to maintain log of HashRing operations
- KnowledgeLogs can be used to avoid full list scans

https://sites.cs.ucsb.edu/~ckrintz/papers/ic2e22.pdf

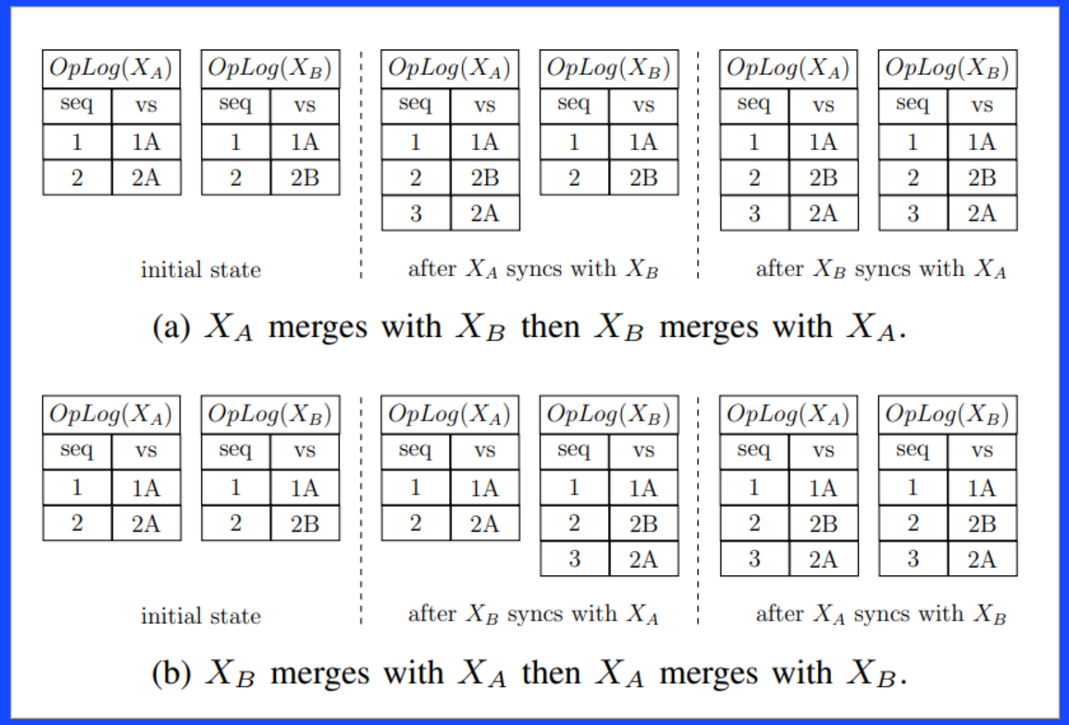

$OpLog(X_A)$

seq	vs
1	1A
2	2A

$OpLog(X_B)$

seq	vs
1	1A
2	2B

initial state

$OpLog(X_A)$

seq	vs
1	1A
2	2B
3	2A

$OpLog(X_B)$

seq	vs
1	1A
2	2B

after X_A syncs with X_B

$OpLog(X_A)$

seq	vs
1	1A
2	2B
3	2A

$OpLog(X_B)$

seq	vs
1	1A
2	2B
3	2A

after X_B syncs with X_A

(a) X_A merges with X_B then X_B merges with X_A.

$OpLog(X_A)$

seq	vs
1	1A
2	2A

$OpLog(X_B)$

seq	vs
1	1A
2	2B

initial state

$OpLog(X_A)$

seq	vs
1	1A
2	2A

$OpLog(X_B)$

seq	vs
1	1A
2	2B
3	2A

after X_B syncs with X_A

$OpLog(X_A)$

seq	vs
1	1A
2	2B
3	2A

$OpLog(X_B)$

seq	vs
1	1A
2	2B
3	2A

after X_A syncs with X_B

(b) X_B merges with X_A then X_A merges with X_B.

CRDTs

- Dot Context
- Dotted Value
- CCounter (uses: Dotted Value)
- MVRegister (uses: DotContext)
- AWSet (uses: Dot Context and Dotted Value)
- AWMap (uses: Dot Context, Dotted Value and AWSet)
- ShoppingListItem (uses: CCounter and MVRegister)
- ShoppingList (uses: AWMap, Dotted Value and ShoppingListItem)

Removing an element from the list - Counter behaviour

- When an item is removed and updated at the same time, it is not removed and its counter is updated instead
- We use an HashMap<string, DottedValue> (integer DV) as a tombstone to indicate the value of the counter when an item is removed.

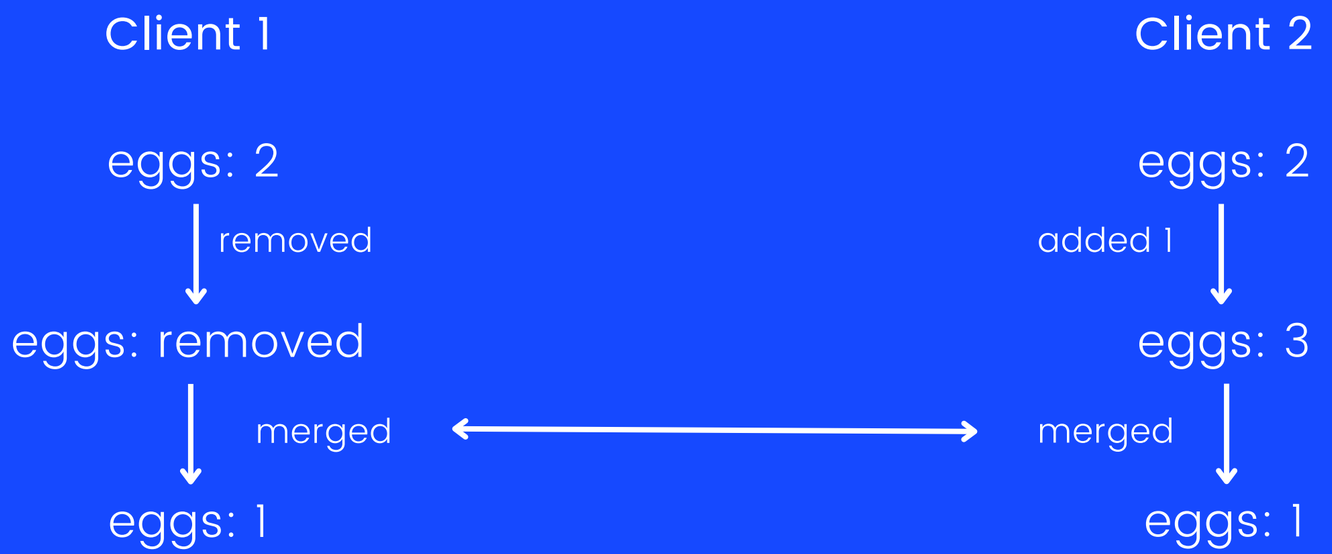

Implementation of the CRDT propagation

- PUT and polling GET HTTP requests

- Since this is a local-first app, the client does not need server approval for making changes, storing it locally first and then sync it via the CrdtSyncService if it is online.

References

- https://www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf
- https://sites.cs.ucsb.edu/~ckrintz/papers/ic2e22.pdf
- https://repositorium.sdum.uminho.pt/bitstream/1822/51503/1/Problem-Solution-Counters-PAPOC2016.pdf